

# softmax\_regression\_2

September 10, 2021

**Name:** Omar Khaled Mahmoud Safwat

**Group:** Alex. G3

```
[ ]: import numpy as np
import pandas as pd

from sklearn import datasets
from sklearn.model_selection import StratifiedShuffleSplit

np.random.seed(42)
```

```
[ ]: iris = datasets.load_iris()
X = iris["data"]
y = iris["target"]
df = pd.DataFrame({fname: values for fname, values in
    ↪zip(iris["feature_names"], X.T)})
df["target"] = y

df.head()
```

```
[ ]:  sepal length (cm)  sepal width (cm)  petal length (cm)  petal width (cm)  \
0                5.1             3.5             1.4             0.2
1                4.9             3.0             1.4             0.2
2                4.7             3.2             1.3             0.2
3                4.6             3.1             1.5             0.2
4                5.0             3.6             1.4             0.2

    target
0        0
1        0
2        0
3        0
4        0
```

## 1 Mini-Batch GD

## 2 Softmax Regression

```
[ ]: class Softmax():  
    """Class Implements Softmax Regression for multinomial classification"""  
  
    def __init__(self):  
        self.X_train = None # User input data.  
        self.y = None # Target feature  
        self.y_ohe = None # Target Variable One-hot encoded  
        self.weights = None # weights of model  
  
    def encode_target(self, y):  
        """One-hot encode target variable"""  
  
        assert y.ndim == 1, "Target variable should be 1 dimensional"  
  
        y_ohe = np.zeros((len(y), y.max() + 1))  
        y_ohe[np.arange(len(y)), y] = 1  
        return y_ohe  
  
    def hypothesis(self, X=None, weights=None, batch_indices=None):  
        """Hypothesis function"""  
        if X is None:  
            X = self.X  
        if weights is None:  
            weights = self.weights  
        if batch_indices is None:  
            batch_indices = list(range(len(self.y)))  
  
        score_fun = X[batch_indices, :] @ weights # Score function for each  
→ class, for each sample  
        score_exp = np.exp(score_fun)  
  
        # Returns a 2D Numpy array, prediction of each class for each sample  
        # Shape: (m, k)  
        # m: Number of samples, k: Number of classes  
        return (score_exp / np.sum(score_exp, axis=1, keepdims=True))  
  
    def loss(self, h, batch_indices):  
        """Cost function"""  
        y_ohe = self.y_ohe[batch_indices, :]  
        h += 1e-7 # Add tolerance term for the log  
        return -np.sum(y_ohe * np.log(h))  
  
    def loss_prime(self, h, batch_indices):  
        """Jacobian vector of cost function"""  
        y_ohe = self.y_ohe[batch_indices, :]
```

```

X = self.X[batch_indices, :]

return (X.T @ (h - y_oh))

def initialize(self, learn_rate, batch_size):
    """
    Initialize first epoch

    Args:
        learn_rate (float): Learning rate for gradient descent

        guess (int, 2D numpy array): Initial guess for model weights

        batch_size (int): Training set batch size
    """
    # Initialize guesses
    self.weights_hist = np.random.randn(self.X.shape[1], self.y_oh.
→shape[1], 1) # Initialize weight vector, account for bias column

    # Initialize history
    self.loss_hist = np.empty(shape=(1))
    self.grad_hist = np.empty_like(self.weights_hist)

    self.n_points = batch_size
    self.epoch = 0
    self.learn_rate = learn_rate

def update_weights_GD(self, idx_1, idx_2, weights):
    """
    Function updates weights along the direction of steepest descent and
→stores results

    Args:
        idx_1 (int): Start index for the current batch

        idx_2 (int): End index for the current batch (non inclusive)

        weights (2D numpy array): Model weights
    """

    # Compute predictions
    self.h_pred = self.hypothesis(batch_indices=self.train_idx[idx_1:
→idx_2], weights=weights)

    # Compute weights gradient
    grad_new = self.loss_prime(self.h_pred, self.train_idx[idx_1: idx_2])
    self.grad_hist = np.dstack((self.grad_hist, np.atleast_3d(grad_new)))

```

```

    # Update Weights
    weights = weights - self.learn_rate * grad_new
    self.weights_hist = np.dstack((self.weights_hist, np.
→atleast_3d(weights)))

    def mini_batch_GD(self, learn_rate=0.01, n_batches=8, max_epochs=1e3,
→seed=None):
        """
        Optimize weights using Mini_batch Gradient Descent

        Args:

            learn_rate (float, optional): Learning rate for gradient descent

            n_batches (int, optional): Number of batches for mini-batch
→Gradient descent

            max_epochs (int, optional): Maximum number of iterations over the
→entire training set

            seed (int, optional): seed the the data shuffling after before
→each epoch

        Returns:

            2d numpy array: Trained model weights
        """
        # Randomly shuffle the dataset's order
        if (seed is not None):
            np.random.seed(seed)

        self.n_batches = n_batches
        np.random.shuffle(self.train_indx)
        self.MAX_EPOCHS = max_epochs

        batch_size = len(self.train_indx) // n_batches

        # Initialize first epoch
        self.initialize(learn_rate, batch_size)

        # Train model using Mini_batches
        while (self.epoch < self.MAX_EPOCHS):
            for i in range(n_batches):
                idx_1 = i * batch_size
                idx_2 = idx_1 + batch_size

```

```

        self.update_weights_GD(idx_1, idx_2, self.weights_hist[:, :, -1]
        ↪-1])

        # Update Predictions and Loss from last epoch
        self.h_pred = self.hypothesis(batch_indices=self.train_indx[:
        ↪batch_size], weights=self.weights_hist[:, :, -1])
        self.loss_hist = np.append(self.loss_hist, self.loss(self.h_pred,
        ↪self.train_indx[:batch_size]))

        # Increment epoch and shuffle data incides for next epoch
        self.epoch += 1
        np.random.shuffle(self.train_indx)

        # Record minimum loss so far if early stopping is used
        if self.early_stop == True:
            # Check if stop criteria threshold is met
            if (self.epoch - self.best_epoch) > self.early_stop_thresh:
                self.is_converged = True
                return self.weights_hist[:, :, -1]

        val_preds = self.hypothesis(batch_indices=self.valid_indx,
        ↪weights=self.weights_hist[:, :, -1])
        current_loss = self.loss(val_preds, self.valid_indx)
        if self.min_val_loss > current_loss:
            self.min_val_loss = current_loss
            self.best_epoch = self.epoch

    return self.weights_hist[:, :, -1]

def fit(self, X, y, early_stop=False, early_stop_thresh=5, **kwargs):
    """Fit model to training data"""

    # X, and y are 2D Numpy arrays.
    # Add a column of ones for theta_0
    self.X = np.hstack((np.ones((len(X), 1)), X))
    self.y = y

    # One-hot encode target variable
    self.y_ohe = self.encode_target(y)
    all_indices = np.array([i for i in range(len(self.X))])

    self.early_stop = early_stop

    # Split the data to validation set to check early stop
    if self.early_stop == True:

```

```

        validation_mask = np.array([np.random.choice([0, 1], p=[0.7, 0.3])
→for i in range(len(y))], dtype='bool')
        self.valid_indx = all_indices[validation_mask]
        self.train_indx = all_indices[~validation_mask]

        self.min_val_loss = float("inf")           # Value of minimum loss
        self.best_epoch = 0                       # Epoch with minimum loss
        self.early_stop_thresh = early_stop_thresh # Maximum number of
→epochs to perform without improvement in loss
        self.is_converged = False

    else:
        self.train_indx = all_indices

    self.weights = self.mini_batch_GD(**kwargs)

def predict(self, X):
    """Returns predictions"""
    X = np.hstack((np.ones((X.shape[0], 1)), X))
    proba = self.hypothesis(X, self.weights, list(range(X.shape[0])))
    return np.argmax(proba, axis=1)

def predict_proba(self, X):
    """Returns predictions for each class as probabilities"""
    X = np.hstack((np.ones((X.shape[0], 1)), X))
    return self.hypothesis(X, self.weights, list(range(X.shape[0])))

def accuracy_score(self, X_test, y_test):
    """Returns model accuracy score"""
    return np.sum(self.predict(X_test) == y_test) / len(y_test)

def show_summary(self):
    """Prints a brief summary after stop criteria is reached"""
    print("Solver summary:")
    print("=" * len("Solver summary:"))

    if self.early_stop == True:
        print("Number of epochs: ", self.best_epoch)
        print("Negative Log Likelihood: ", self.loss_hist[self.best_epoch])
        print("Early stop criteria was reached first: ", self.is_converged)
        print("Train accuracy: ", self.accuracy_score(self.X[self.
→train_indx, 1:], self.y[self.train_indx]))
        print("Validation accuracy: ", self.accuracy_score(self.X[self.
→valid_indx, 1:], self.y[self.valid_indx]))
    else:
        print("Number of epochs: ", self.epoch)
        print("Negative Log Likelihood: ", self.loss_hist[-1])

```

```
print("Train accuracy: ", self.accuracy_score(self.X[self.
↪train_indx, 1:], self.y[self.train_indx]))
```

### 3 Train with Mini\_batch GD

#### 3.1 Without Early stop

```
[ ]: # Split dataset into validation and training
all_indices = np.array([i for i in range(len(X))])
np.random.shuffle(all_indices)
validation_mask = np.array([np.random.choice([0, 1], p=[0.7, 0.3]) for i in
↪range(len(y))], dtype='bool')
train_indx, val_indx = all_indices[~validation_mask],
↪all_indices[validation_mask]

X_train, y_train = X[train_indx], y[train_indx]
X_val, y_val = X[val_indx], y[val_indx]

softmax = Softmax()
softmax.fit(X_train, y_train, early_stop=False, learn_rate=0.05)
softmax.show_summary()
print('Validation accuracy: ', softmax.accuracy_score(X_val, y_val))
```

Solver summary:

=====

Number of epochs: 1000  
 Negative Log Likelihood: 0.5375553478186633  
 Train accuracy: 0.9904761904761905  
 Validation accuracy: 0.9777777777777777

#### 3.2 With Early stop

```
[ ]: softmax = Softmax()
softmax.fit(X, y, early_stop=True, early_stop_thresh=300, learn_rate=0.05)
softmax.show_summary()
```

Solver summary:

=====

Number of epochs: 66  
 Negative Log Likelihood: 0.37147882094226636  
 Early stop criteria was reached first: True  
 Train accuracy: 0.979381443298969  
 Validation accuracy: 0.9622641509433962

## 4 Cross Validation

```
[ ]: class KfoldCV():
    """Class implements k-fold cross validation"""

    def __init__(self, model):
        self.model = model # ML model object
        self.accuracy_scores=None

    def fit(self, X, y, n_folds=3, **kwargs):
        """Train machine learning model using k_folds cross validation"""

        self.accuracy_scores = []

        # Split Dataset
        split = StratifiedShuffleSplit(n_splits=n_folds, test_size=0.2,
        random_state=42)
        for train_index, test_index in split.split(X, y):
            X_train, y_train = X[train_index], y[train_index]
            X_test, y_test = X[test_index], y[test_index]

            self.model.fit(X_train, y_train, **kwargs)
            self.accuracy_scores.append(self.model.accuracy_score(X_test, y_test))

    def predict(self, X):
        """Return model predictions for dataset X"""
        return self.model.predict(X)

    def predict_proba(self, X):
        """Return model probability predictions for dataset X"""
        return self.model.predict_proba(X)

    def show_summary(self):
        self.model.show_summary()

    def accuracy_score(self):
        """Returns average accuracy score from CV"""
        return np.mean(self.accuracy_scores)
```

### 4.1 CV without early stopping

```
[ ]: softmaxCV = KfoldCV(Softmax())
softmaxCV.fit(X, y, learn_rate=0.05, early_stop=False)
print("Cross Validation accuracy: ", softmaxCV.accuracy_score())
```

Cross Validation accuracy: 0.9111111111111111



## 4.2 CV with early stopping

```
[ ]: softmaxCV = KfoldCV(Softmax())  
softmaxCV.fit(X, y, learn_rate=0.05, early_stop=True, early_stop_thresh=300) #  
    ↪ With early stop criteria  
print("Cross Validation accuracy: ", softmaxCV.accuracy_score())
```

Cross Validation accuracy: 0.9444444444444443